

Problem 3

```
library(e1071)
library(ISLR)
attach(OJ)
```

a)

```
set.seed(2017)

train = sample(1:nrow(OJ), 535)
test = setdiff(1:nrow(OJ), train)
```

b)

```
svm.fit <- svm(Purchase~.,data=OJ[train,],kernel="linear",cost=1)
summary(svm.fit)

##
## Call:
## svm(formula = Purchase ~ ., data = OJ[train, ], kernel = "linear",
##      cost = 1)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel: linear
##              cost: 1
##
## Number of Support Vectors:  202
##
## ( 102 100 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

The SVM machine has 202 support vectors with 102 and 100 in the classes.

c) What are the training and test error rates?

```
yhat <- predict(svm.fit)

table(predict=yhat, truth=OJ$Purchase[train])

##      truth
## predict  CH  MM
##      CH 293  40
##      MM  37 165
```

The training error rate is the count of errors made over the tootal count or $(40+37)/(201+80+132+122) = 0.14$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##      truth
## predict CH  MM
##      CH 201 132
##      MM 122  80
```

The test error rate is the count of errors made over the count of cases or $(132+122)/(201+80+132+122) = 0.47$.

d)

```
tune.out <- tune(svm,Purchase~.,data=OJ[train,],kernel="linear",ranges=list(cost=c(0.01,0.05,0.1,0.5,0.7,0.85,1,5,7.5,8,10)),
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.5
##
## - best performance: 0.1531447
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.1626485 0.04984485
## 2  0.05 0.1607268 0.07040721
## 3  0.10 0.1607268 0.06642371
## 4  0.50 0.1531447 0.06044108
## 5  0.70 0.1549965 0.05928271
## 6  0.75 0.1549965 0.05928271
## 7  0.85 0.1549965 0.05928271
## 8  0.95 0.1549965 0.05928271
## 9  1.00 0.1549965 0.05928271
## 10 5.00 0.1643606 0.05659160
## 11 7.50 0.1662823 0.06116933
## 12 8.00 0.1681342 0.05957297
## 13 10.00 0.1681342 0.05695706
```

An optimal cost is 0.75.

e)

```
bestmod <- tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = OJ[train,
##           ], ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 0.7, 0.75, 0.85,
##           0.95, 1, 5, 7.5, 8, 10)), kernel = "linear")
##
```

```
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  0.5
##
## Number of Support Vectors:  207
##
## ( 103 104 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

```
yhat <- predict(bestmod)

table(predict=yhat, truth=OJ$Purchase[train])
```

```
##          truth
## predict  CH  MM
##       CH 294  39
##       MM  36 166
```

The training error rate is the count of errors made over the total count or $(40+36)/535 = 0.14$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##          truth
## predict  CH  MM
##       CH 201 132
##       MM 122  80
```

The test error rate is the count of errors made over the tootal count or $(133+122)/535 = 0.48$

- (f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma.

```
svmrاد.fit <- svm(Purchase~.,data=OJ[train,],kernel="radial")
summary(svmrad.fit)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ[train, ], kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  1
##
## Number of Support Vectors:  260
##
## ( 127 133 )
##
##
## Number of Classes:  2
```

```
##
## Levels:
## CH MM
```

The SVM machine has 260 support vectors with 127 and 133 in the classes.

```
yhat <- predict(svmrad.fit)

table(predict=yhat, truth=OJ$Purchase[train])
```

```
##          truth
## predict CH  MM
##      CH 306  49
##      MM  24 156
```

The training error rate is the count of errors made over the total count or $(49+24)/535 = 0.136$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##          truth
## predict CH  MM
##      CH 220 135
##      MM 103  77
```

The test error rate is the count of errors made over the total count or $(135+103)/535 = 0.445$

```
tunerad.out <- tune(svm,Purchase~.,data=OJ[train,],kernel="radial",ranges=list(cost=c(0.01,0.05,0.1,0.5
summary(tunerad.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 5
##
## - best performance: 0.1682739
##
## - Detailed performance results:
## cost      error dispersion
## 1 0.01 0.3829490 0.04681606
## 2 0.05 0.3232006 0.04590639
## 3 0.10 0.2017470 0.04827853
## 4 0.50 0.1831936 0.05047356
## 5 1.00 0.1776380 0.04857852
## 6 5.00 0.1682739 0.05086527
## 7 10.00 0.1701607 0.04657436
```

Best model is at a cost of 5.

```
bestmodrad <- tunerad.out$best.model
summary(bestmodrad)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = OJ[train,
##           ], ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 1, 5, 10)), kernel = "radial")
```

```
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  5
##
## Number of Support Vectors:  218
##
## ( 105 113 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

This SVM has 218 support vectors with 105 and 113 in each class.

```
yhat <- predict(bestmodrad)
```

```
table(predict=yhat, truth=OJ$Purchase[train])
```

```
##           truth
## predict  CH  MM
##       CH 306  42
##       MM  24 163
```

The training error rate is the count of errors made over the total count or $(42+24)/535 = 0.123$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##           truth
## predict  CH  MM
##       CH 216 132
##       MM 107  80
```

The test error rate is the count of errors made over the total count or $(132+107)/535 = 0.447$

g)

```
svmply.fit <- svm(Purchase~.,data=OJ[train,],kernel="polynomial",degree=2)
summary(svmply.fit)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ[train, ], kernel = "polynomial",
##     degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##       cost:  1
##     degree:  2
##    coef.0:  0
##
```

```
## Number of Support Vectors: 315
##
## ( 153 162 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

The SVM now has 315 support vectors with 153 and 162 in the classes.

```
yhat <- predict(svmplay.fit)

table(predict=yhat, truth=OJ$Purchase[train])
```

```
##      truth
## predict CH  MM
##      CH 310  80
##      MM  20 125
```

The training error rate is the count of errors made over the total count or $(80+20)/535 = 0.187$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##      truth
## predict CH  MM
##      CH 238 152
##      MM  85  60
```

The test error rate is the count of errors made over the total count or $(152+85)/535 = 0.443$

```
tuneply.out <- tune(svm,Purchase~.,data=OJ[train,],kernel="polynomial",ranges=list(cost=c(0.01,0.05,0.1
summary(tuneply.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 10
##
## - best performance: 0.1722572
##
## - Detailed performance results:
## cost      error dispersion
## 1  0.01 0.3630678 0.08121140
## 2  0.05 0.3405311 0.06305619
## 3  0.10 0.3199161 0.06900220
## 4  0.50 0.1984277 0.05271605
## 5  1.00 0.1946191 0.05263791
## 6  5.00 0.1816212 0.04802431
## 7 10.00 0.1722572 0.04681925
```

The best model is at a cost of 10.

```
bestmodply <- tuneply.out$best.model
summary(bestmodply)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = OJ[train,
##           ], ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 1, 5, 10)), kernel = "polynomial")
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:    10
##   degree:    3
##   coef.0:    0
##
## Number of Support Vectors:  215
##
##   ( 108 107 )
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

This SVM has 215 support vectors with 108 and 108 in the classes.

```
yhat <- predict(bestmodply)

table(predict=yhat, truth=OJ$Purchase[train])
```

```
##           truth
## predict  CH  MM
##      CH 313  43
##      MM  17 162
```

The training error rate is the count of errors made over the total count or $(43+17)/535 = 0.112$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##           truth
## predict  CH  MM
##      CH 221 135
##      MM 102  77
```

The test error rate is the count of errors made over the total count or $(135+102)/535 = 0.443$

h)

The results are not as good. This model is simpler in that it is a linear model; only the predictors have been transformed. This is a similar approach to transforming the predictors for a linear regression model.

```
svmexp.fit <- svm(Purchase~.+poly(PriceDiff,2)+poly(ListPriceDiff,2)+poly(DiscMM,2)+poly(LoyalCH,2)+poly(LoyalMM,2))
summary(svmexp.fit)
```

```
##
## Call:
```

```
## svm(formula = Purchase ~ . + poly(PriceDiff, 2) + poly(ListPriceDiff,
##      2) + poly(DiscMM, 2) + poly(LoyalCH, 2) + poly(SalePriceMM, 2) +
##      poly(SalePriceCH, 2) + poly(PctDiscMM, 2) + poly(PctDiscCH, 2),
##      data = OJ[train, ], kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1
##
## Number of Support Vectors:  200
##
## ( 97 103 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

This SVM has 200 support vectors with 97 and 103 in the classes.

```
yhat <- predict(svmexp.fit)
```

```
table(predict=yhat, truth=OJ$Purchase[train])
```

```
##      truth
## predict CH  MM
##      CH 295  43
##      MM  35 162
```

The training error rate is the count of errors made over the total count or $(43+34)/535 = 0.146$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##      truth
## predict CH  MM
##      CH 206 132
##      MM 117  80
```

The test error rate is the count of errors made over the total count or $(132+117)/535 = 0.465$

```
tuneexp.out <- tune(svm,Purchase~.+poly(PriceDiff,2)+poly(ListPriceDiff,2)+poly(DiscMM,2)+poly(LoyalCH,2)+poly(SalePriceMM,2)+poly(SalePriceCH,2)+poly(PctDiscMM,2)+poly(PctDiscCH,2),
summary(tuneexp.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     5
##
## - best performance: 0.1701957
##
## - Detailed performance results:
```



```
##      cost      error dispersion
## 1  0.01 0.1737945 0.03041157
## 2  0.05 0.1776380 0.04102085
## 3  0.10 0.1720475 0.04234664
## 4  0.50 0.1739693 0.04862185
## 5  1.00 0.1758211 0.04625692
## 6  5.00 0.1701957 0.04552464
## 7 10.00 0.1720825 0.04251418
```

```
bestmodexp <- tuneexp.out$best.model
summary(bestmodexp)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ . + poly(PriceDiff,
##      2) + poly(ListPriceDiff, 2) + poly(DiscMM, 2) + poly(LoyalCH,
##      2) + poly(SalePriceMM, 2) + poly(SalePriceCH, 2) + poly(PctDiscMM,
##      2) + poly(PctDiscCH, 2), data = OJ[train, ], ranges = list(cost = c(0.01,
##      0.05, 0.1, 0.5, 1, 5, 10)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  5
##
## Number of Support Vectors:  197
##
## ( 98 99 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

This optimal SVM has a cost of 10 and 195 support vectors with 97 and 98 in the 2 classes.

```
yhat <- predict(bestmodexp)

table(predict=yhat, truth=OJ$Purchase[train])
```

```
##      truth
## predict CH  MM
##      CH 297  43
##      MM  33 162
```

The training error rate is the count of errors made over the total count or $(42+36)/535 = 0.146$

```
table(predict=yhat, truth=OJ$Purchase[test])
```

```
##      truth
## predict CH  MM
##      CH 208 132
##      MM 115  80
```

The test error rate is the count of errors made over the total count or $(130+117)/535 = 0.462$

i)

The polynomial model has the best fit for the test data (0.443). Its 215 support vectors indicate that it is also at the mid-range of complexity for the various SVMs assessed here.